

Computer Lab 1 (732A99 Machine Learning)

Lennart Schilling (lensc874)

22 November 2018

Assignment 1

1.1

At first, the data from the Excel file *spambase.xlsx* will be imported and splitted into train and test data (50%:50%)

```
# Importing data
library(readxl)
data = read_excel("spambase.xlsx")

# Dividing data into train and test set
n = dim(data)[1]
set.seed(12345)
id = sample(1:n, floor(n*0.5))
train = data[id,]
test = data[-id,]
```

1.2

Using the train data, a logistic regression model will be created. Analysing the p-values of the coefficients, it can be seen which independent variables have a significant influence on the dependent variable *Spam*. For the sake of clarity, this is not explicitly stated in this report.

```
# Fitting model
logitModel = suppressWarnings(glm(Spam ~ ., data = train, family = binomial))
summary(logitModel)
```

In the next step, the *logitModel* will be used to classify emails of the training and test data. To prevent duplicate code in 1.3, the *classificationLogit*-function was coded. Giving data and a threshold as an input, a list with the specified threshold to decide which probabilities lead to a spam classification, the confusion matrix and the misclassification rate will be returned.

```
# Classifying & evaluating results
classificationLogit = function(data, threshold = 0.5) {
  # Classifying emails with the model
  yFit = predict(logitModel,
    newdata = data[, !colnames(data) %in% "Spam"],
    type='response')
  yFit = ifelse(yFit > threshold, 1, 0)
  # Evaluating classification results
  confusionMatrix = table(y = data$Spam, yFit)
  misclassificationRate <- mean(yFit != data$Spam)
  # Returning results
  return(
    list(
      threshold = threshold,
      confusionMatrix = confusionMatrix,
      misclassificationRate = misclassificationRate
    )
  )
}
```

```
)
)
}
```

Using the train data and the default threshold (0.5) as the input leads to the following confusion matrix and misclassification rate.

```
classificationLogitTrain = classificationLogit(data = train)
classificationLogitTrain$confusionMatrix
```

```
##      yFit
## y      0    1
## 0 803 142
## 1  81 344
```

```
classificationLogitTrain$misclassificationRate
```

```
## [1] 0.1627737
```

Instead, using the test data and the default threshold (0.5) as the input leads to the following confusion matrix and misclassification rate.

```
classificationLogitTest = classificationLogit(data = test)
classificationLogitTest$confusionMatrix
```

```
##      yFit
## y      0    1
## 0 791 146
## 1  97 336
```

```
classificationLogitTest$misclassificationRate
```

```
## [1] 0.1773723
```

It can be seen that the model performs about equally well for both data. The misclassification rate is slightly better for the training data (16.2%) than for the training data (17.7%). This is an indication that there is not too much overfitting on the training data.

1.3

Now we are changing the classification principle and therefore the input threshold from 0.5 to 0.9.

Using the train data and the threshold = 0.9 as the input leads to the following confusion matrix and misclassification rate.

```
classificationLogitTrainAdjThreshold = classificationLogit(data = train, threshold = 0.9)
classificationLogitTrainAdjThreshold$confusionMatrix
```

```
##      yFit
## y      0    1
## 0 944  1
## 1 419  6
```

```
classificationLogitTrainAdjThreshold$misclassificationRate
```

```
## [1] 0.3065693
```

Using the test data and the threshold = 0.9 as the input leads to the following confusion matrix and misclassification rate.

```
classificationLogitTestAdjThreshold = classificationLogit(data = test, threshold = 0.9)
classificationLogitTestAdjThreshold$confusionMatrix
```

```
##      yFit
## y      0   1
## 0 936   1
## 1 427   6
```

```
classificationLogitTestAdjThreshold$misclassificationRate
```

```
## [1] 0.3124088
```

In both cases it can be seen that the classification quality decreases a lot (misclassification rates about 30%). Because the threshold is now much higher than before, the number of false negative predictions has increased strongly.

1.4

In the following, the standard classifier *kknn()* was used to predict spam mails. Again, to prevent duplicate code in 1.5, the *classificationKknn*-function was coded. Giving data, number of k and a threshold as an input, a list with the same elements as the *classificationLogit*-function is returned.

```
library(kknn)
# Classifying & evaluating results
classificationKknn = function(data, k, threshold = 0.5) {
  # Classifying emails
  kknnModel <- kknn(formula = Spam ~ .,
                    train = train,
                    test = data,
                    k = k)
  kknnModel$fitted.values = ifelse(kknnModel$fitted.values > threshold, 1, 0)
  # Evaluating classification results
  confusionMatrix = table(y = data$Spam, yFit = kknnModel$fitted.values)
  misclassificationRate <- mean(kknnModel$fitted.values != data$Spam)
  # Returning results
  return(
    list(
      threshold = threshold,
      confusionMatrix = confusionMatrix,
      misclassificationRate = misclassificationRate
    )
  )
}
```

Using the train data and $k = 30$ as the input leads to the following misclassification rate.

```
classificationKnnTrain = classificationKknn(data = train, k = 30)
classificationKnnTrain$misclassificationRate
```

```
## [1] 0.1722628
```

Instead, using the test data and $k = 30$ as the input leads to the following misclassification rate.

```
classificationKnnTest = classificationKknn(data = test, k = 30)
classificationKnnTest$misclassificationRate
```

```
## [1] 0.329927
```

Here, a big difference between the prediction power of the model related to the train data (misclassification rate: 17.2%) and the test data (misclassification rate: 32.9%) can be observed. This leads to the assumption that the model is overfitting on the training data. Compared to the results of the logistic regression model with the threshold = 0.5, this model does not deliver such accurate predictions.

1.5

Now we are changing the k from 30 to 1.

Using the train data and k = 1 as the input leads to the following misclassification rate.

```
classificationKnnTrain = classificationKknn(data = train, k = 1)
classificationKnnTrain$misclassificationRate
```

```
## [1] 0
```

Using the test data and k = 1 as the input leads to the following misclassification rate.

```
classificationKnnTest = classificationKknn(data = test, k = 1)
classificationKnnTest$misclassificationRate
```

```
## [1] 0.3459854
```

This example shows very clearly that the model is strongly overfitted on the training data. While it classifies every mail for the training data correctly, the misclassification rate for the test data is almost 35%. With k = 1, the classification depends only on the nearest neighbor (the value of the dependent variable of this observation in the training data where the independent variables have the lowest distance to the observation which shall be classified) which leads to a much higher dependency on the training data.

Assignment 3

3.1

```
featureSelection = function(X, Y, Nfolds) {
  # Checking input
  if (!is.matrix(X)) {
    stop("X must be of class 'matrix'.")
  }
  if (!is.numeric(Y)) {
    stop("Y must be a numeric vector.")
  }
  if (!is.numeric(Nfolds) | length(Nfolds) != 1) {
    stop("Nfolds must be a numeric of length 1.")
  }
  if (nrow(X) != length(Y)) {
    stop("X and Y must have same number of observations.")
  }
  # Adding intercept variable as last column to X
  X = cbind(X, Intercept = 1)
  # Creating folds
  set.seed(12345)
  obs = 1:nrow(X)
  folds = list()
  # Allocating observations equally to folds
  for (i in 1:Nfolds) {
    folds[[i]] = sample(obs, floor(nrow(X)/Nfolds))
  }
}
```

```

    obs = obs[!obs %in% folds[[i]]]
  }
  # Creating equal sized folds could lead to unallocated observations
  # which will be allocated one by one iteratively
  i = 1
  while (length(obs) > 0) {
    folds[[i]] = c(folds[[i]], obs[1])
    obs = obs[-1]
    i = i+1
    if (i > Nfolds) {
      i = 1
    }
  }
}
# Identifying all possible subsets
i = 1
allSubsets = list()
for (nX in 0:(ncol(X)-1)) {
  combinations = combn(1:(ncol(X)-1), nX)
  for (comb in 1:ncol(combinations)) {
    allSubsets[[i]] = c(combinations[,comb], ncol(X)) # to each subset the intercept will be also add
    i = i+1
  }
}
# Selecting best subset for linear regression
# Setting up list object with CV-score (default = 0) for each subset
subsetCVScore = list()
for (i in 1:length(allSubsets)) {
  subsetCVScore[[i]] = 0
}
i = 1
# Calculating CV-score for each subset applying k-cross-validation
for (subset in allSubsets) {
  subsetX = as.matrix(X[,subset])
  for (validationSet in folds) {
    # Creating train and validation set
    trainX = subsetX[-validationSet,]
    trainY = Y[-validationSet]
    validationX = subsetX[validationSet,]
    validationY = Y[validationSet]
    # Getting beta estimates with train data
    betaEstimates = (solve(t(trainX) %*% trainX)) %*% t(trainX) %*% trainY
    # Estimating validationY
    validationYFit = validationX %*% betaEstimates
    # Calculating MSE
    mse = mean((validationY - validationYFit)^2)
    # Adding MSE to subsetCVScore
    subsetCVScore[[i]] = subsetCVScore[[i]] + mse
  }
  # calculating CV-score for each subset
  subsetCVScore[[i]] = subsetCVScore[[i]]/Nfolds
  i = i+1
}
# Identifying subset with least CV-score

```

```

leastCVScore = 99999999
for (i in 1:length(subsetCVScore)) {
  if (subsetCVScore[[i]] < leastCVScore) {
    leastCVScore = subsetCVScore[[i]]
    bestSubset = i
  }
}
# Calculating data frame with CV-score per number of features for plot
nFeaturesCVScore = setNames(data.frame(matrix(ncol = 2, nrow = 0)), c("nFeatures", "cvScore"))
i = 1
for (subset in allSubsets) {
  nFeaturesCVScore[i,1] = length(allSubsets[[i]])
  nFeaturesCVScore[i,2] = subsetCVScore[[i]]
  i = i+1
}
# Returning elements
print("Best subset:")
print(sort(colnames(X[,allSubsets[[bestSubset]]])))
print(paste("CV score: ", round(subsetCVScore[[bestSubset]], 2), sep = ""))
ggplot2::ggplot(nFeaturesCVScore, ggplot2::aes(nFeatures, cvScore)) +
  ggplot2::geom_point() +
  ggplot2::theme_bw() +
  ggplot2::ggtitle("CV-score vs. Number of features (Intercept included)")
}

```

3.2

```

data(swiss)
featureSelection(X = as.matrix(swiss[,-1]),
  Y = swiss[,1],
  Nfolds = 5)


```

```

## [1] "Best subset:"
## [1] "Agriculture"      "Catholic"          "Education"
## [4] "Infant.Mortality" "Intercept"
## [1] "CV score: 56.51"

```

CV-score vs. Number of features (Intercept included)



It can be seen that with an increasing number of features, the CV score approximately decreases linearly. While the CV score is very high using only the *intercept* as a predictor ($nFeatures = 1$), the minimum CV score of 56.51 can be found for $nFeatures = 5$ (which are the variables *Agriculture*, *Catholic*, *Education*, *Infant.Mortality* and the *intercept*).

Linear regression identified a linear relationship between the independent variables and the dependent variable *Fertility*. To say that it is reasonable that the independent variables have largest impact on the target, causation needs to be checked. It indicates that one event is the result of the occurrence of the other event.


Assignment 4

4.1

```
# Importing data
library(readxl)
data = read_excel("tecator.xlsx")

# Plotting moisture versus protein
library(ggplot2)
ggplot(data = data, aes(x = Protein, y = Moisture)) +
  geom_point() +
  geom_smooth(method = "lm") +
  theme_bw() +
  labs(title = "Moisture vs. Protein")
```

Moisture vs. Protein



Since the blue line indicates a linear model which assumes that there is a positive relationship between Moisture and Protein, the data can be presented by a linear model. However, the model would not predict the outlier values which can be observed in the top left corner of the plot.

4.2

The probabilistic model for M_i can be described as

$$M_i = \beta_0 + \beta_1 Protein + \beta_2 Protein^2 + \dots + \beta_i Protein^i + \epsilon.$$

Since we know that Moisture is normally distributed and

$$\epsilon \sim N(0, \sigma^2),$$

$$M_i \sim N\left(\sum_{j=0}^i \beta_j Protein^j, \sigma^2\right).$$

Since

$$\epsilon = \sum_{j=0}^i (y_j - \hat{y}_j)$$

and

$$MSE = \frac{1}{n} \sum_{j=0}^i (y_j - \hat{y}_j)^2,$$

it becomes clear that the calculation of the MSE is very similar to the calculation of the error term. That shows that the MSE is a good indicator for how well the model fits the data, because it quantifies the difference between the predicted and actual values.

4.3

```
# Dividing data into train and test set
n = dim(data)[1]
set.seed(12345)
id = sample(1:n, floor(n*0.5))
train = data[id,]
test = data[-id,]

# Fitting models
m1 = lm(Moisture ~ Protein, data = train)
m2 = lm(Moisture ~ Protein + I(Protein^2), data = train)
m3 = lm(Moisture ~ Protein + I(Protein^2) + I(Protein^3), data = train)
m4 = lm(Moisture ~ Protein + I(Protein^2) + I(Protein^3) + I(Protein^4), data = train)
m5 = lm(Moisture ~ Protein + I(Protein^2) + I(Protein^3) + I(Protein^4) + I(Protein^5), data = train)
m6 = lm(Moisture ~ Protein + I(Protein^2) + I(Protein^3) + I(Protein^4) + I(Protein^5) + I(Protein^6), data = train)

# Predicting moisture values using the fitted models
# Train data
m1yFitTrain = predict(m1, train)
m2yFitTrain = predict(m2, train)
m3yFitTrain = predict(m3, train)
m4yFitTrain = predict(m4, train)
m5yFitTrain = predict(m5, train)
m6yFitTrain = predict(m6, train)
# Test data
m1yFitTest = predict(m1, test)
m2yFitTest = predict(m2, test)
m3yFitTest = predict(m3, test)
m4yFitTest = predict(m4, test)
m5yFitTest = predict(m5, test)
m6yFitTest = predict(m6, test)

# Calculating MSE
# Train data
m1MSETrain = mean((train$Moisture - m1yFitTrain)^2)
m2MSETrain = mean((train$Moisture - m2yFitTrain)^2)
m3MSETrain = mean((train$Moisture - m3yFitTrain)^2)
m4MSETrain = mean((train$Moisture - m4yFitTrain)^2)
m5MSETrain = mean((train$Moisture - m5yFitTrain)^2)
m6MSETrain = mean((train$Moisture - m6yFitTrain)^2)
# Test data
m1MSETest = mean((test$Moisture - m1yFitTest)^2)
m2MSETest = mean((test$Moisture - m2yFitTest)^2)
m3MSETest = mean((test$Moisture - m3yFitTest)^2)
m4MSETest = mean((test$Moisture - m4yFitTest)^2)
m5MSETest = mean((test$Moisture - m5yFitTest)^2)
m6MSETest = mean((test$Moisture - m6yFitTest)^2)


# Combining results in a data frame
finalDf = setNames(data.frame(matrix(ncol = 3, nrow = 12)), c("i", "data", "mse"))
finalDf$i = c(1:6, 1:6)
finalDf$data = c(rep("train", 6), rep("test", 6))
finalDf$mse = c(m1MSETrain, m2MSETrain, m3MSETrain, m4MSETrain, m5MSETrain, m6MSETrain,
```

```

m1MSETest, m2MSETest, m3MSETest, m4MSETest, m5MSETest, m6MSETest)

# Plotting results
ggplot(data = finalDf, aes(x = i, y = mse, color = data)) +
  geom_point() +
  theme_bw() +
  labs(title = "MSE vs. i")

```



Which model is best according to the plot? How do the MSE values change and why? Interpret this picture in terms of bias-variance tradeoff.

Since the least difference between the test and train MSE can be observed for $i=1$, model M_1 is the best according to the plot. The test MSE is even lower than the training MSE, which even indicates little underfitting. A big challenge in fitting a model is to avoid overfitting. As model flexibility increases, training MSE will decrease, but the test MSE may increase as shown in the plot as well. This happens if the statistical model identifies patterns in the training data which are just caused by random chance.

4.4

```

library(MASS)
model = lm(Fat ~ ., data = data[,2:102])
stepwiseSelection = stepAIC(model, direction = "both", trace = FALSE)
summary(stepwiseSelection)

##
## Call:
## lm(formula = Fat ~ Channel1 + Channel2 + Channel4 + Channel5 +

```

```

##      Channel7 + Channel8 + Channel11 + Channel12 + Channel13 +
##      Channel14 + Channel15 + Channel17 + Channel19 + Channel20 +
##      Channel22 + Channel24 + Channel25 + Channel26 + Channel28 +
##      Channel29 + Channel30 + Channel32 + Channel34 + Channel36 +
##      Channel37 + Channel39 + Channel40 + Channel41 + Channel42 +
##      Channel45 + Channel46 + Channel47 + Channel48 + Channel50 +
##      Channel51 + Channel52 + Channel54 + Channel55 + Channel56 +
##      Channel59 + Channel60 + Channel61 + Channel63 + Channel64 +
##      Channel65 + Channel67 + Channel68 + Channel69 + Channel71 +
##      Channel73 + Channel74 + Channel78 + Channel79 + Channel80 +
##      Channel81 + Channel84 + Channel85 + Channel87 + Channel88 +
##      Channel92 + Channel94 + Channel98 + Channel99, data = data[,
##      2:102])
##
## Residuals:
##      Min        1Q      Median        3Q        Max
## -2.82961 -0.57129 -0.00696  0.58152  2.86375
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      7.093      1.453   4.882 2.64e-06 ***
## Channel11     10559.894    2333.430   4.525 1.21e-05 ***
## Channel12     -12636.967    3467.995  -3.644 0.000369 ***
## Channel14       8489.323    4637.993   1.830 0.069164 .
## Channel15     -10408.967    4771.350  -2.182 0.030689 *
## Channel17      -5376.018    3851.782  -1.396 0.164847
## Channel18       7215.595    4246.489   1.699 0.091342 .
## Channel111     -9505.520    5721.115  -1.661 0.098692 .
## Channel112      37240.918   12290.648   3.030 0.002878 **
## Channel113     -41564.547   15892.375  -2.615 0.009817 **
## Channel114      34938.179   13290.454   2.629 0.009454 **
## Channel115     -23761.451    6584.006  -3.609 0.000417 ***
## Channel117       4296.572    3189.730   1.347 0.179998
## Channel119      14279.808    5017.407   2.846 0.005042 **
## Channel120     -23855.616    5153.161  -4.629 7.85e-06 ***
## Channel122      18444.906    3381.683   5.454 1.97e-07 ***
## Channel124     -20138.426    4946.417  -4.071 7.52e-05 ***
## Channel125      18137.432    5374.094   3.375 0.000938 ***
## Channel126     -7670.318    3859.006  -1.988 0.048660 *
## Channel128      20079.898    4991.631   4.023 9.06e-05 ***
## Channel129     -36351.014    7655.223  -4.749 4.72e-06 ***
## Channel130      18071.276    5863.802   3.082 0.002446 **
## Channel132       3838.013    2722.862   1.410 0.160729
## Channel134     -9242.884    2225.926  -4.152 5.48e-05 ***
## Channel136       8070.938    3317.588   2.433 0.016152 *
## Channel137     -9045.588    3536.621  -2.558 0.011522 *
## Channel139      18664.454    5986.730   3.118 0.002183 **
## Channel140     -20069.709   10701.902  -1.875 0.062677 .
## Channel141      22257.776   11122.533   2.001 0.047169 *
## Channel142     -21760.853    5833.811  -3.730 0.000270 ***
## Channel145      18145.804    2985.416   6.078 9.50e-09 ***
## Channel146     -8225.696    3715.367  -2.214 0.028330 *
## Channel147     -4986.549    2558.694  -1.949 0.053165 .
## Channel148       2876.075    2014.985   1.427 0.155546

```

```

## Channel50    -13009.410    4535.797    -2.868  0.004720 **
## Channel51     29251.161    6554.297     4.463  1.57e-05 ***
## Channel52    -26833.976    4389.473    -6.113  7.97e-09 ***
## Channel54     30954.862    4392.339     7.047  6.06e-11 ***
## Channel55    -35183.287    5646.314    -6.231  4.39e-09 ***
## Channel56     14912.986    2810.889     5.305  3.93e-07 ***
## Channel59     -8030.278    1887.431    -4.255  3.66e-05 ***
## Channel60     13071.416    2629.374     4.971  1.79e-06 ***
## Channel61     -7850.189    2246.864    -3.494  0.000625 ***
## Channel63     15059.275    3231.692     4.660  6.90e-06 ***
## Channel64    -19909.466    4727.696    -4.211  4.35e-05 ***
## Channel65       4190.184    3486.766     1.202  0.231346
## Channel67     13850.508    3909.121     3.543  0.000526 ***
## Channel68    -25873.365    5304.223    -4.878  2.69e-06 ***
## Channel69     18362.385    3331.483     5.512  1.50e-07 ***
## Channel71     -9223.910    1558.752    -5.917  2.11e-08 ***
## Channel73     12456.498    2386.255     5.220  5.82e-07 ***
## Channel74     -5624.411    1933.590    -2.909  0.004177 **
## Channel78     -7927.105    2176.860    -3.642  0.000372 ***
## Channel79     15473.188    3812.200     4.059  7.89e-05 ***
## Channel80    -22391.895    4490.714    -4.986  1.67e-06 ***
## Channel81     13852.453    3105.934     4.460  1.59e-05 ***
## Channel84    -11442.630    3457.064    -3.310  0.001167 **
## Channel85     20228.671    4081.863     4.956  1.91e-06 ***
## Channel87    -15938.315    4102.273    -3.885  0.000153 ***
## Channel88       5647.072    3236.286     1.745  0.083033 .
## Channel92       6595.995    1864.595     3.537  0.000537 ***
## Channel94     -5497.846    1847.113    -2.976  0.003397 **
## Channel98     -8728.596    2489.314    -3.506  0.000598 ***
## Channel99       8554.587    1898.010     4.507  1.31e-05 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.107 on 151 degrees of freedom
## Multiple R-squared:  0.9947, Adjusted R-squared:  0.9925
## F-statistic: 447.9 on 63 and 151 DF,  p-value: < 2.2e-16

nVariables = length(stepwiseSelection$coefficients)-1
paste0("selected variables: ", nVariables)

## [1] "selected variables: 63"

```

4.5

```

library(glmnet)

## Loading required package: Matrix
## Loading required package: foreach
## Loaded glmnet 2.0-16


covariates = data[,2:101]
response = data[,102]
modelRidge = glmnet(as.matrix(covariates),
                    as.matrix(response),

```

```

alpha=0,
family="gaussian")
plot(modelRidge, xvar="lambda", label=TRUE)

```




All coefficients converge to zero as Log Lambda increases.

4.6

```

modelLasso = glmnet(as.matrix(covariates),
                    as.matrix(response),
                    alpha=1,
                    family="gaussian")
plot(modelLasso, xvar = "lambda", label = TRUE)

```



Compared to the Ridge regression, using Lasso regression, the coefficients converge much faster towards zero as Log Lambda increases.

4.7

```
modelLassoCV = cv.glmnet(as.matrix(covariates),
                        as.matrix(response),
                        alpha = 1,
                        family = "gaussian",
                        lambda = seq(0,1,0.001))
coef(modelLassoCV, s = "lambda.min")
```

```
## 101 x 1 sparse Matrix of class "dgCMatrix"
##              1
## (Intercept)  1.218582e+01
## Channel1    2.243351e+01
## Channel2    -5.477212e+01
## Channel3     2.028288e+01
## Channel4     1.790966e+01
## Channel5     1.599792e+01
## Channel6     1.496901e+01
## Channel7     4.930552e+01
## Channel8     1.410886e+01
## Channel9     1.100549e+01
## Channel10    -3.946699e+01
## Channel11    -3.413441e+01
## Channel12    -3.361613e+00
```

```

## Channel13      1.501427e+02
## Channel14      2.372668e-01
## Channel15     -4.867299e+01
## Channel16     -4.549011e+01
## Channel17     -3.891421e+01
## Channel18     -3.095813e+01
## Channel19     -2.488455e+01
## Channel20     -2.152392e+01
## Channel21     -1.721212e+01
## Channel22     -1.288916e+01
## Channel23     -7.467292e+00
## Channel24      3.858146e-01
## Channel25      6.141436e+00
## Channel26      7.339759e+00
## Channel27      4.952231e+00
## Channel28      9.256704e-01
## Channel29     -1.845498e+00
## Channel30     -6.201764e+00
## Channel31     -1.240428e+01
## Channel32     -1.608200e+01
## Channel33     -1.533188e+01
## Channel34     -8.878726e+00
## Channel35     -1.222622e+00
## Channel36      4.986875e+00
## Channel37      8.663342e+00
## Channel38      9.507388e+00
## Channel39      5.835387e+00
## Channel40      5.486958e+00
## Channel41      1.437055e+02
## Channel42      2.782707e+01
## Channel43      9.118839e+00
## Channel44     -9.479487e+00
## Channel45     -2.436273e+00
## Channel46     -7.205962e+00
## Channel47     -7.485645e-04
## Channel48     -1.494150e+01
## Channel49     -1.826387e+01
## Channel50     -2.915222e+01
## Channel51     -6.977177e+01
## Channel52     -1.219494e+01
## Channel53      7.071090e+00
## Channel54     -9.397099e+00
## Channel55     -3.634639e-01
## Channel56      4.870685e+00
## Channel57      3.943612e+01
## Channel58      1.141937e+01
## Channel59      4.812258e+00
## Channel60      5.489478e+00
## Channel61      1.251985e+01
## Channel62      2.021891e+01
## Channel63     -7.405961e+00
## Channel64     -2.464353e+01
## Channel65      3.321800e+01
## Channel66     -9.834681e+00

```

```

## Channel67      1.262018e+01
## Channel68     -3.217675e+00
## Channel69      8.903740e+00
## Channel70      9.503287e-01
## Channel71      3.949972e+00
## Channel72      1.789541e+01
## Channel73     -2.043606e+01
## Channel74     -1.276925e+01
## Channel75     -8.831294e+00
## Channel76     -2.240484e+01
## Channel77      8.138153e-01
## Channel78     -2.734175e+01
## Channel79     -2.300507e+01
## Channel80      2.379017e-04
## Channel81      1.391527e-04
## Channel82      9.866988e-05
## Channel83      1.303538e-06
## Channel84     -8.303980e-05
## Channel85     -4.297569e-05
## Channel86     -3.530336e-05
## Channel87      5.470530e-05
## Channel88      8.946494e-05
## Channel89      9.648548e-05
## Channel90     -2.640082e+01
## Channel91      6.485393e-01
## Channel92      7.039651e-01
## Channel93     -1.993822e+01
## Channel94      2.455602e+01
## Channel95      2.755495e+01
## Channel96      4.056484e-02
## Channel97      8.631612e+00
## Channel98      4.537254e+00
## Channel99      5.168688e+00
## Channel100     1.640315e+01

paste0("Selected variables: ", length(coef(modelLassoCV, s = "lambda.min"))-1) # all variables used


## [1] "Selected variables: 100"

paste0("Optimal Lambda: ", modelLassoCV$lambda.min)

## [1] "Optimal Lambda: 0"

plot(modelLassoCV)

```

4.8

Comparing the two different methods, the number of variables selected differs strongly. While as a result of 4.4, 63 variables were selected, in 4.7 all 100 independent variables were selected to predict the dependent variable.