

Computer lab 1 block 2: Ensemble methods mixture models (732A99 Machine Learning)

Lennart Schilling (*lensc874*)

03 December 2018

Contents

1. Ensemble Methods	1
2. Mixture Models	3
2.1 Code	3
2.2 K=2	6
2.3 K=3	12
2.4 K=4	28
2.5 Conclusion	54

1. Ensemble Methods

At first, the data from the Excel file *spambase.csv* will be imported and splitted into train and test data.

```
library(randomForest)
library(mboost)
library(ggplot2)

# importing data
data = read.csv2("spambase.csv")

# converting 'Spam' to 'factor' so that randomForest() does classification instead of regression
data$Spam = as.factor(data$Spam)

# dividing data into train and test set
n = dim(data)[1]
set.seed(12345)
id = sample(1:n, floor(n*(2/3)))
train = data[id,]
test = data[-id,]
```

To evaluate the performance of Adaboost classification trees and random forests on the spam data, the function *modelComparison* automatically creates a plot which compares the error rates of the two algorithms for the in the input specified different number of trees (*maxNTree*). Both models will be trained with the specified train set (*trainData*). The error rates are related to the specified test set (*testX* & *testY*).

```
modelComparison = function(formula, trainData, testX, testY, maxNTree) {
  # calculating errors for random forest
  allErrors = as.data.frame(cbind(nTrees = seq(from = 10, to = maxNTree, by = 10),
                                   errorRate = randomForest(formula = formula,
                                                             data = trainData,
                                                             xtest = testX,
                                                             ytest = testY,
                                                             ntree = maxNTree)$test$err.rate[
```

```

seq(from = 10,
    to = maxNTree,
    by = 10),1],

    model = "randomForest"))
# calculating errors for adaBoost
for (i in seq(from = 10, to = maxNTree, by = 10)) {
  adaBoostModel = blackboost(formula = formula,
                             data = trainData,
                             family = AdaExp(),
                             control = boost_control(mstop = i))
  yFit = predict(object = adaBoostModel,
                 newdata = testX,
                 type = "class")
  error = 1 - sum(ifelse(as.numeric(as.character(testY)) ==
                        as.numeric(as.character(yFit)), 1, 0)) / length(testY)
  allErrors = rbind(allErrors,
                    cbind(nTrees = i, errorRate = error, model = "adaBoost"))
}
# adjusting classes
allErrors$nTrees = as.numeric(as.character(allErrors$nTrees))
allErrors$errorRate = as.numeric(as.character(allErrors$errorRate))
allErrors$model = as.character(allErrors$model)
# plotting data
ggplot(data = allErrors,
       mapping = aes(x = nTrees, y = errorRate, color = model)) +
  geom_point(size = 2) +
  theme_bw() +
  ggtitle("Error rate vs. number of trees")
}

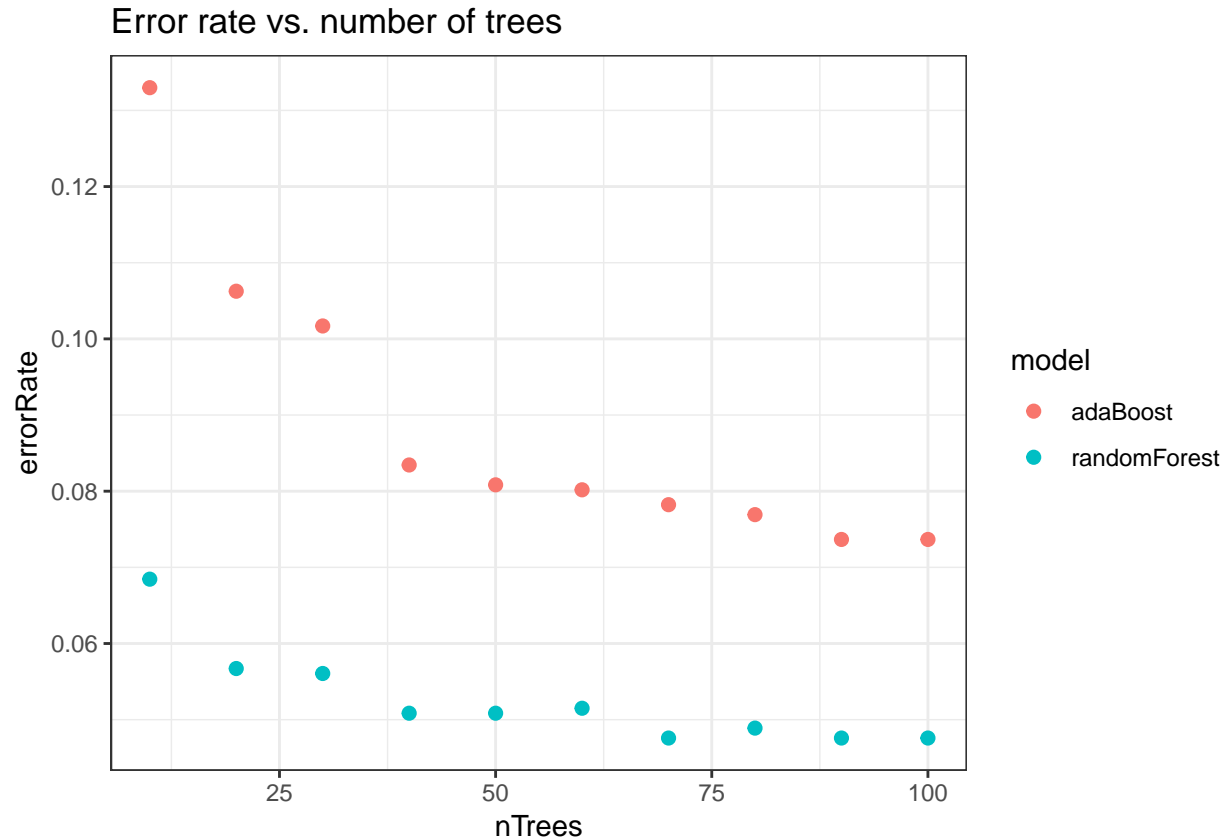
```

The following output of the function shows that the random forest algorithm generally leads to a lower test error rate for this data. In this case, the optimal choice of trees is 50, because the test error rate is minimal. Instead, the ADA boosted algorithm shows higher error rates. Especially for a small number of trees the algorithm does not deliver satisfying results compared to the random forest algorithm.

```

modelComparison(formula = Spam ~ .,
                 trainData = train,
                 testX = test[,!colnames(data) %in% "Spam"],
                 testY = test$Spam,
                 maxNTree = 100)

```



2. Mixture Models

2.1 Code

To compare the results for $K = 2, 3, 4$, the *em*-function provides a graphical analysis for every iteration. The function includes comments which explain what I did at which step to create the EM algorithm. The function will be finally run with $K = 2, 3, 4$.

```
em = function(K) {
  # Initializing data
  set.seed(1234567890)
  max_it = 100 # max number of EM iterations
  min_change = 0.1 # min change in log likelihood between two consecutive EM iterations
  N = 1000 # number of training points
  D = 10 # number of dimensions
  x = matrix(nrow=N, ncol = D) # training data
  true_pi = vector(length = K) # true mixing coefficients
  true_mu = matrix(nrow = K, ncol = D) # true conditional distributions
  true_pi = c(rep(1/K, K))
  if (K == 2) {
    true_mu[1,] = c(0.5,0.6,0.4,0.7,0.3,0.8,0.2,0.9,0.1,1)
    true_mu[2,] = c(0.5,0.4,0.6,0.3,0.7,0.2,0.8,0.1,0.9,0)
    plot(true_mu[1,], type = "o", xlab = "dimension", col = "blue",
         ylim = c(0,1), main = "True")
    points(true_mu[2,], type="o", xlab = "dimension", col = "red",
           main = "True")
  }
}
```

```

} else if (K == 3) {
  true_mu[1,] = c(0.5,0.6,0.4,0.7,0.3,0.8,0.2,0.9,0.1,1)
  true_mu[2,] = c(0.5,0.4,0.6,0.3,0.7,0.2,0.8,0.1,0.9,0)
  true_mu[3,] = c(0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5)
  plot(true_mu[1,], type = "o", xlab = "dimension", col = "blue", ylim=c(0,1),
        main = "True")
  points(true_mu[2,], type = "o", xlab = "dimension", col = "red",
          main = "True")
  points(true_mu[3,], type = "o", xlab = "dimension", col = "green",
          main = "True")
} else {
  true_mu[1,] = c(0.5,0.6,0.4,0.7,0.3,0.8,0.2,0.9,0.1,1)
  true_mu[2,] = c(0.5,0.4,0.6,0.3,0.7,0.2,0.8,0.1,0.9,0)
  true_mu[3,] = c(0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5)
  true_mu[4,] = c(0.3,0.5,0.5,0.7,0.5,0.5,0.5,0.5,0.4,0.5)
  plot(true_mu[1,], type = "o", xlab = "dimension", col = "blue",
        ylim = c(0,1), main = "True")
  points(true_mu[2,], type = "o", xlab = "dimension", col = "red",
          main = "True")
  points(true_mu[3,], type = "o", xlab = "dimension", col = "green",
          main = "True")
  points(true_mu[4,], type = "o", xlab = "dimension", col = "yellow",
          main = "True")
}

z = matrix(nrow = N, ncol = K) # fractional component assignments
pi = vector(length = K) # mixing coefficients
mu = matrix(nrow = K, ncol = D) # conditional distributions
llik = vector(length = max_it) # log likelihood of the EM iterations
# Producing the training data
for(n in 1:N) {
  k = sample(1:K, 1, prob = true_pi)
  for(d in 1:D) {
    x[n,d] = rbinom(1, 1, true_mu[k,d])
  }
}

# Random initialization of the paramters
pi = runif(K, 0.49, 0.51)
pi = pi / sum(pi)
for(k in 1:K) {
  mu[k,] = runif(D, 0.49, 0.51)
}

# EM algorithm
for(it in 1:max_it) {
  # Plotting mu
  # Defining plot title
  title = paste0("Iteration", it)
  if (K == 2) {
    plot(mu[1,], type = "o", xlab = "dimension", col = "blue", ylim = c(0,1), main = title)
    points(mu[2,], type = "o", xlab = "dimension", col = "red", main = title)
  } else if (K == 3) {
    plot(mu[1,], type = "o", xlab = "dimension", col = "blue", ylim = c(0,1), main = title)
    points(mu[2,], type = "o", xlab = "dimension", col = "red", main = title)
    points(mu[3,], type = "o", xlab = "dimension", col = "green", main = title)
  }
}

```

```

} else {
  plot(mu[1,], type = "o", xlab = "dimension", col = "blue", ylim = c(0,1), main = title)
  points(mu[2,], type = "o", xlab = "dimension", col = "red", main = title)
  points(mu[3,], type = "o", xlab = "dimension", col = "green", main = title)
  points(mu[4,], type = "o", xlab = "dimension", col = "yellow", main = title)
}
Sys.sleep(0.5)
# E-step: Computation of the fractional component assignments
for (n in 1:N) {
  # Creating empty matrix (column 1:K = p_x_given_k; column K+1 = p(x/all k)
  p_x = matrix(data = c(rep(1,K), 0), nrow = 1, ncol = K+1)
  # Calculating p(x/k) and p(x/all k)
  for (k in 1:K) {
    # Calculating p(x/k)
    for (d in 1:D) {
      p_x[1,k] = p_x[1,k] * (mu[k,d]^x[n,d]) * (1-mu[k,d])^(1-x[n,d])
    }
    p_x[1,k] = p_x[1,k] * pi[k] # weighting with pi[k]
    # Calculating p(x/all k) (denominator)
    p_x[1,K+1] = p_x[1,K+1] + p_x[1,k]
  }
  # Calculating z for n and all k
  for (k in 1:K) {
    z[n,k] = p_x[1,k] / p_x[1,K+1]
  }
}
# Log likelihood computation
for (n in 1:N) {
  for (k in 1:K) {
    log_term = 0
    for (d in 1:D) {
      log_term = log_term + x[n,d] * log(mu[k,d]) + (1-x[n,d]) * log(1-mu[k,d])
    }
    llik[it] = llik[it] + z[n,k] * (log(pi[k]) + log_term)
  }
}
cat("iteration: ", it, "log likelihood: ", llik[it], "\n")
flush.console()
# Stop if the log likelihood has not changed significantly
if (it != 1) {
  if (abs(llik[it] - llik[it-1]) < min_change) {
    break
  }
}
# M-step: ML parameter estimation from the data and fractional component assignments
# Updating pi
for (k in 1:K) {
  pi[k] = sum(z[,k])/N
}
# Updating mu
for (k in 1:K) {
  mu[k,] = 0
  for (n in 1:N) {

```

```


    mu[k,] = mu[k,] + x[n,] * z[n,k]
  }
  mu[k,] = mu[k,] / sum(z[,k])
}
}
# Printing pi, mu and development of log likelihood at the end
return(list(
  pi = pi,
  mu = mu,
  logLikelihoodDevelopment = plot(llik[1:it],
                                  type = "o",
                                  main = "Development of the log likelihood",
                                  xlab = "iteration",
                                  ylab = "log likelihood")
))
}

```


2.2 K=2

First, the function will be run for $K=2$.


```
em(2)
```




```
## iteration: 1 log likelihood: -7623.897
```




```
## iteration: 2 log likelihood: -7610.745
```




iteration: 3 log likelihood: -7463.445




iteration: 4 log likelihood: -6575.121




iteration: 5 log likelihood: -5731.559




iteration: 6 log likelihood: -5656.174




iteration: 7 log likelihood: -5648.904




iteration: 8 log likelihood: -5646.139




iteration: 9 log likelihood: -5644.608




iteration: 10 log likelihood: -5643.615




iteration: 11 log likelihood: -5642.913




iteration: 12 log likelihood: -5642.386




iteration: 13 log likelihood: -5641.977




iteration: 14 log likelihood: -5641.649




iteration: 15 log likelihood: -5641.382




iteration: 16 log likelihood: -5641.161




iteration: 17 log likelihood: -5640.975




iteration: 18 log likelihood: -5640.819




iteration: 19 log likelihood: -5640.685



iteration: 20 log likelihood: -5640.571



iteration: 21 log likelihood: -5640.473




\$pi
[1] 0.5110531 0.4889469
##

```
## $mu
##      [,1]      [,2]      [,3]      [,4]      [,5]      [,6]      [,7]
## [1,] 0.4931735 0.3974606 0.5967811 0.2785480 0.6927917 0.2184957 0.8018491
## [2,] 0.4989543 0.6255823 0.3804363 0.7171478 0.3230343 0.7778699 0.2049559
##      [,8]      [,9]      [,10]
## [1,] 0.1116477 0.88054439 0.004290353
## [2,] 0.9140913 0.08997919 0.999714736
##
## $logLikelihoodDevelopment
## NULL
```


2.3 K=3

Next, the function will be run for $K=3$.


```
em(3)
```




```
## iteration: 1 log likelihood: -8029.723
```




```
## iteration: 2 log likelihood: -8027.183
```




iteration: 3 log likelihood: -8024.696




iteration: 4 log likelihood: -8005.631




iteration: 5 log likelihood: -7877.606




iteration: 6 log likelihood: -7403.513




iteration: 7 log likelihood: -6936.919




iteration: 8 log likelihood: -6818.582




iteration: 9 log likelihood: -6791.377




iteration: 10 log likelihood: -6780.713




iteration: 11 log likelihood: -6774.958




iteration: 12 log likelihood: -6771.261




iteration: 13 log likelihood: -6768.606




iteration: 14 log likelihood: -6766.535




iteration: 15 log likelihood: -6764.815




iteration: 16 log likelihood: -6763.316




iteration: 17 log likelihood: -6761.967




iteration: 18 log likelihood: -6760.727




iteration: 19 log likelihood: -6759.572




iteration: 20 log likelihood: -6758.491




iteration: 21 log likelihood: -6757.475




iteration: 22 log likelihood: -6756.521




iteration: 23 log likelihood: -6755.625




iteration: 24 log likelihood: -6754.784




iteration: 25 log likelihood: -6753.996




iteration: 26 log likelihood: -6753.26




iteration: 27 log likelihood: -6752.571




iteration: 28 log likelihood: -6751.928




iteration: 29 log likelihood: -6751.328




iteration: 30 log likelihood: -6750.768




iteration: 31 log likelihood: -6750.246




iteration: 32 log likelihood: -6749.758




iteration: 33 log likelihood: -6749.304




iteration: 34 log likelihood: -6748.88




iteration: 35 log likelihood: -6748.484




iteration: 36 log likelihood: -6748.114




iteration: 37 log likelihood: -6747.767




iteration: 38 log likelihood: -6747.444




iteration: 39 log likelihood: -6747.14




iteration: 40 log likelihood: -6746.856




iteration: 41 log likelihood: -6746.589




iteration: 42 log likelihood: -6746.338




iteration: 43 log likelihood: -6746.102




iteration: 44 log likelihood: -6745.88




iteration: 45 log likelihood: -6745.67




iteration: 46 log likelihood: -6745.472




iteration: 47 log likelihood: -6745.285




iteration: 48 log likelihood: -6745.108




iteration: 49 log likelihood: -6744.939




iteration: 50 log likelihood: -6744.78




iteration: 51 log likelihood: -6744.627




iteration: 52 log likelihood: -6744.483




iteration: 53 log likelihood: -6744.344




iteration: 54 log likelihood: -6744.212




iteration: 55 log likelihood: -6744.086




iteration: 56 log likelihood: -6743.964




iteration: 57 log likelihood: -6743.848




iteration: 58 log likelihood: -6743.736




iteration: 59 log likelihood: -6743.628




```
## iteration: 60 log likelihood: -6743.524
```



```
## iteration: 61 log likelihood: -6743.423
```



```
## iteration: 62 log likelihood: -6743.326
```




```
## $pi
## [1] 0.3259592 0.3044579 0.3695828
##
```

```
## $mu
##      [,1]      [,2]      [,3]      [,4]      [,5]      [,6]      [,7]
## [1,] 0.4737193 0.3817120 0.6288021 0.3086143 0.6943731 0.1980896 0.7879447
## [2,] 0.4909874 0.4793213 0.4691560 0.4791793 0.5329895 0.4928830 0.4643990
## [3,] 0.5089571 0.5834802 0.4199272 0.7157107 0.2905703 0.7667258 0.2320784
##      [,8]      [,9]      [,10]
## [1,] 0.1349651 0.8912534 0.01937869
## [2,] 0.4902682 0.4922194 0.39798407
## [3,] 0.8516111 0.1072226 0.99981353
##
## $logLikelihoodDevelopment
## NULL
```


2.4 K=4

Finally, the function will be run for $K=4$.


```
em(4)
```




```
## iteration: 1 log likelihood: -8316.904
```




```
## iteration: 2 log likelihood: -8291.114
```




iteration: 3 log likelihood: -8286.966




iteration: 4 log likelihood: -8264.806




iteration: 5 log likelihood: -8161.19




iteration: 6 log likelihood: -7868.89




iteration: 7 log likelihood: -7570.873




iteration: 8 log likelihood: -7445.719




iteration: 9 log likelihood: -7389.741




iteration: 10 log likelihood: -7356.803




iteration: 11 log likelihood: -7337.208




iteration: 12 log likelihood: -7326.118




iteration: 13 log likelihood: -7319.998




iteration: 14 log likelihood: -7316.6




iteration: 15 log likelihood: -7314.666




iteration: 16 log likelihood: -7313.528




iteration: 17 log likelihood: -7312.829




iteration: 18 log likelihood: -7312.367




iteration: 19 log likelihood: -7312.024




iteration: 20 log likelihood: -7311.723




iteration: 21 log likelihood: -7311.407




iteration: 22 log likelihood: -7311.036




iteration: 23 log likelihood: -7310.574




iteration: 24 log likelihood: -7309.988




iteration: 25 log likelihood: -7309.248




iteration: 26 log likelihood: -7308.322




iteration: 27 log likelihood: -7307.185




iteration: 28 log likelihood: -7305.809




iteration: 29 log likelihood: -7304.176




iteration: 30 log likelihood: -7302.273




iteration: 31 log likelihood: -7300.1




iteration: 32 log likelihood: -7297.671




iteration: 33 log likelihood: -7295.014




iteration: 34 log likelihood: -7292.171




iteration: 35 log likelihood: -7289.196




iteration: 36 log likelihood: -7286.15




iteration: 37 log likelihood: -7283.093




iteration: 38 log likelihood: -7280.079




iteration: 39 log likelihood: -7277.151




iteration: 40 log likelihood: -7274.34




iteration: 41 log likelihood: -7271.66




iteration: 42 log likelihood: -7269.116




iteration: 43 log likelihood: -7266.7




iteration: 44 log likelihood: -7264.398




iteration: 45 log likelihood: -7262.189




iteration: 46 log likelihood: -7260.051




iteration: 47 log likelihood: -7257.96




iteration: 48 log likelihood: -7255.892




iteration: 49 log likelihood: -7253.824




iteration: 50 log likelihood: -7251.733




iteration: 51 log likelihood: -7249.603




iteration: 52 log likelihood: -7247.419




iteration: 53 log likelihood: -7245.17




iteration: 54 log likelihood: -7242.853




iteration: 55 log likelihood: -7240.472




iteration: 56 log likelihood: -7238.038




iteration: 57 log likelihood: -7235.571




iteration: 58 log likelihood: -7233.095




iteration: 59 log likelihood: -7230.64




iteration: 60 log likelihood: -7228.239




iteration: 61 log likelihood: -7225.925




iteration: 62 log likelihood: -7223.725




iteration: 63 log likelihood: -7221.663




iteration: 64 log likelihood: -7219.755




iteration: 65 log likelihood: -7218.01




iteration: 66 log likelihood: -7216.431




iteration: 67 log likelihood: -7215.013




iteration: 68 log likelihood: -7213.748




iteration: 69 log likelihood: -7212.621




iteration: 70 log likelihood: -7211.62




iteration: 71 log likelihood: -7210.727




iteration: 72 log likelihood: -7209.929




iteration: 73 log likelihood: -7209.208




iteration: 74 log likelihood: -7208.552




iteration: 75 log likelihood: -7207.946




iteration: 76 log likelihood: -7207.38




iteration: 77 log likelihood: -7206.844




iteration: 78 log likelihood: -7206.327




iteration: 79 log likelihood: -7205.824




iteration: 80 log likelihood: -7205.326




iteration: 81 log likelihood: -7204.829




iteration: 82 log likelihood: -7204.327




iteration: 83 log likelihood: -7203.816




iteration: 84 log likelihood: -7203.294




iteration: 85 log likelihood: -7202.756




iteration: 86 log likelihood: -7202.201




iteration: 87 log likelihood: -7201.627




iteration: 88 log likelihood: -7201.032




iteration: 89 log likelihood: -7200.414




iteration: 90 log likelihood: -7199.773




iteration: 91 log likelihood: -7199.107




iteration: 92 log likelihood: -7198.416




iteration: 93 log likelihood: -7197.7




iteration: 94 log likelihood: -7196.957




iteration: 95 log likelihood: -7196.188




iteration: 96 log likelihood: -7195.392




iteration: 97 log likelihood: -7194.57




iteration: 98 log likelihood: -7193.722



```
## iteration: 99 log likelihood: -7192.847
```



```
## iteration: 100 log likelihood: -7191.946
```



```
## $pi
## [1] 0.2880470 0.2533761 0.2933710 0.1652060
##
## $mu
##      [,1]      [,2]      [,3]      [,4]      [,5]      [,6]      [,7]
## [1,] 0.3714855 0.3899958 0.4790260 0.5731886 0.5022651 0.5108478 0.2835691
## [2,] 0.5199997 0.6135841 0.3891214 0.7132736 0.2722448 0.7785461 0.2168891
## [3,] 0.4383456 0.4042497 0.5489526 0.3298363 0.6578057 0.2049012 0.7825505
## [4,] 0.3428531 0.7784238 0.5591637 0.6319621 0.5167044 0.4629058 0.7311279
##      [,8]      [,9]     [,10]
## [1,] 0.3519184 0.36924863 0.48252239
## [2,] 0.9337959 0.08504806 0.99916297
## [3,] 0.1703330 0.80517853 0.04500171
## [4,] 0.6601375 0.46532151 0.48814639
##
## $logLikelihoodDevelopment
```

NULL

2.5 Conclusion

Comparing the final plots for each of the cases, it becomes clear that when the mixture model has more components ($K = 4$), the EM algorithm does not perform as accurate as for fewer components ($K = 2$ or $K = 3$). The segregation between each component gets diluted as the components get higher.